

Day 9 - White noise and random walks

Introduction

Last week, we introduced the Holt-Winters filter, a model that is capable of estimating a time series with trend and seasonal components. However, sometimes the Holt-Winters model is not able to capture the serial correlation present in the data, resulting in left-over serial correlation in the *residual error series*. In order to capture serial correlation in the residual error series, we need to build more complicated models by combining time series techniques used for the mean and error structure of the series. This week, we begin considering basic stochastic models for time series, including: white noise, random walks, and autoregressive models.

To guide our discussion, we will use a [data set](#) describing the number of visits to a website for a statistics course at Duke University, called Statistical Forecasting. The data set contains daily counts of the number of visits to the website between 2014 and 2020.

```
# packages
library(tidyverse)
library(lubridate)

# read table from the web
web <- read_csv("daily-website-visitors.csv")
```

Pros:

- Filled in notes
- Office hours

Changes:

- More live coding, examples, and activities
- = Randomized groups :)
- Slow it down Benson Boone style

Cons:

- No printed notes :)
- Not finishing notes
- Balancing class, mid-terms, job search, etc

Review


Construct a time series of the daily unique visits to the website in the year of 2017, then fit a Holt-Winters model to the data and plot the resulting model. Then create an ACF plot of the residual error series and comment on what you see.

💡 Tip


You can use the `resid` function on a Holt-Winters model to obtain the residual error series.

```
web2017 <- web %>%  
  mutate(dt = mdy(Date)) %>%  
  filter(year(dt) == 2017) %>%  
  arrange(dt)  
  
web_ts <- ts(  
  web2017$Unique.Visits,  
  start = c(1,1),  
  freq = 7  
)  
  
hw <- HoltWinters(web_ts)  
plot(hw)
```

Holt-Winters filtering



```
acf(resid(hw))
```



i Note

A time series $\{w_t : t = 1, 2, \dots, n\}$ is called discrete white noise (DWN) if the variables w_1, w_2, \dots, w_n are independent and identically distributed with a mean of 0.

If the variables also follow a normal distribution, the series is called

Gaussian noise.

$$w_t \sim N(0, \sigma^2)$$

normal
variance = σ^2
mean = 0

i Second-order properties

$$\mu(t) = 0$$

$$\gamma_k = \text{Cov}(w_t, w_{t+k}) = \begin{cases} \sigma^2 & k = 0 \\ 0 & k \neq 0 \end{cases}$$



$$\rho_k = \begin{cases} 1 & k = 0 \\ 0 & k \neq 0 \end{cases}$$

How can we simulate Gaussian white noise?



```
set.seed(10062024)
par(mfrow = c(3, 2))
for(i in 1:3) {
  tmp <- rnorm(100)
  plot(
    tmp, type = "l",
    xlab = "Time"
  )
  abline(h = 0, lty = 2)
  acf(tmp)
}
```

generate 100 random
normal observations



Series tmp



Series tmp



Series tmp



Random walks

i Note

A time series $\{x_t\}$ is a random walk if

$$x_t = x_{t-1} + w_t$$

where w_t is discrete white noise. It can be shown that







$$x_t = w_t + w_{t-1} + w_{t-2} + \dots$$

i Second-order properties

$$\begin{aligned}\mu(t) &= 0 \\ \gamma_k(t) &= \text{Cov}(x_t, x_{t+k}) = t\sigma^2 \\ \rho_k(t) &= \frac{\text{Cov}(x_t, x_{t+k})}{\sqrt{\text{Var}(x_t)\text{Var}(x_{t+k})}} = \frac{t\sigma^2}{\sqrt{t\sigma^2(t+k)\sigma^2}} = \frac{1}{\sqrt{1+k/t}}\end{aligned}$$

How can we simulate a random walk?

```
set.seed(10062024)
par(mfrow = c(3, 2))
for(i in 1:3) {
  x <- w <- rnorm(500)
  for(t in 2:500) x[t] <- x[t-1] + w[t]
  plot(x, type = "l", xlab = "Time")
  abline(h = 0, lty = 2)
  acf(x)
}
```



```

set.seed(10062024)
par(mfrow = c(2, 2))
for(i in 1:2) {
  # generate random walk
  x <- w <- rnorm(500)
  for(t in 2:500) x[t] <- x[t-1] + w[t]



  # first order difference
  x_dif <- diff(x)
  plot(x_dif, type = "l", xlab = "Time")
  abline(h = 0, lty = 2)
  acf(x_dif)
}

```



first order difference

$$X_t - X_{t-1}$$

First order differences of a random walk are realizations of white noise.



Series x_dif



i Backshift operators

The backward shift operator B is defined by

$$Bx_t = x_{t-1}$$

and

$$B^n x_t = x_{t-n}$$

which provides a convenient way to express autoregressive series.

Example: Express the random walk model in terms of the backshift operator.

$$X_t = X_{t-1} + \omega_t$$

$$X_t = B X_t + \omega_t$$

$$X_t - B X_t = \omega_t$$

$$\underbrace{(1 - B)}_{\text{characteristic polynomial}} X_t = \omega_t$$

↳ characteristic polynomial

Random walk with drift

i Note

A time series $\{x_t\}$ is a random walk w/ drift if


$$x_t = x_{t-1} + \delta + w_t$$

where w_t is white noise.

```
set.seed(10062024)
par(mfrow = c(1,1))
# generate random walk with drift
x <- w <- rnorm(500)
for(t in 2:500) x[t] <- x[t-1] + .10 + w[t]

plot(x, xlab = "Time", type = "l")
```

Handwritten annotations: A circle around `x[t-1]` and a circle around `.10` with an arrow pointing to a handwritten δ .



```
set.seed(10062024)
par(mfrow = c(2, 2))
# generate random walk with drift
x <- w <- rnorm(500)
for(t in 2:500) x[t] <- x[t-1] + .5 + w[t]



# first order difference - notice the mean
x_dif <- diff(x)
mean(x_dif)

[1] 0.5361847
```

Handwritten notes: A circled `.5` in the for loop is pointed to by a line from the text "notice the mean". The output `[1] 0.5361847` is also circled. A handwritten δ is next to the for loop.

```
plot(x_dif, type = "l", xlab = "Time")
abline(h = 0, lty = 2)
acf(x_dif)

# we can fix that
x_dif_centered <- x_dif - mean(x_dif)
plot(x_dif_centered, type = "l", xlab = "Time")
abline(h = 0, lty = 2)
acf(x_dif_centered)
```

Series x_dif**Series x_dif_centered**