

Day 13 - Intro to Regression

Introduction

Trends in time series can be classified as either *stochastic*, *deterministic*, or both. Recently, we have considered *stochastic* models for time series, culminating in the $AR(p)$ process. Stochastic time series models are adept at explaining serial autocorrelation, but can struggle to explain large structural effects, such as trends or seasonality.

To remedy this problem, we often pair stochastic models with models capable of accounting for deterministic trends, such as regression. In today's lab, we will begin to explore ~~time series~~ ^{regression} as a tool for modeling ^{time series}.

Regression overview

Not linear model: $y = e^{\beta_0 + \beta_1 x}$
 $y = \beta_0 + \beta_1 x$

i Linear model theory

A linear model is a model of the form

$$y_i = \underbrace{\beta_0}_{\text{intercept}} + \underbrace{\beta_1 x_{i,1} + \beta_2 x_{i,2} + \dots + \beta_p x_{i,p}}_{\text{identically}} + \underbrace{\epsilon_i}_{\text{error}}$$

where ϵ_i is independent and identically distributed $N(0, \sigma^2)$. In matrix notation, the above model is equivalent to

$$y = X\beta + \epsilon$$

where

$$y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}_{n \times 1}, \quad X = \begin{bmatrix} 1 & x_{11} & x_{12} & \dots & x_{1p} \\ 1 & x_{21} & x_{22} & \dots & x_{2p} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_{n1} & x_{n2} & \dots & x_{np} \end{bmatrix}_{n \times p}, \quad \beta = \begin{bmatrix} \beta_1 \\ \beta_2 \\ \vdots \\ \beta_p \end{bmatrix}_{p \times 1}, \quad \epsilon = \begin{bmatrix} \epsilon_1 \\ \epsilon_2 \\ \vdots \\ \epsilon_n \end{bmatrix}_{n \times 1}$$

Design matrix unknown

By properties of normal distributions, the above model is equivalent to

$$y \sim \mathcal{N}(X\beta, \Sigma)$$

where $\Sigma = \sigma^2 \mathcal{I}_n$ and \mathcal{I}_n is an $n \times n$ identity matrix. To fit the model, we must estimate β 's and σ^2 . The estimated regression equation is written as:

$$\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_{i,1} + \hat{\beta}_2 x_{i,2} + \dots + \hat{\beta}_p x_{i,p}$$

Parameter estimates are obtained by minimizing the sum of squared error. The error in regression is called a residual, (SSE)

minimize: $(y - X\beta)^T (y - X\beta)$ $e_i = y_i - \hat{y}_i$ minimize $\sum e_i^2$

The ordinary least squares estimator (OLS), which minimizes the SSE, is

$$\hat{\beta} = (X^T X)^{-1} X^T y$$

Assumptions

1. Independence of observations.

Tools:

- ACF, PACE (correlogram)
- Context

2. Constant variance

Tools:

- Residuals vs. fitted
- Scale-location plot



3. Linearity: no trends in the residuals

Tools:

- Residuals vs. fitted

4. Normality of errors

Tools:

- QQ plot

5. No influential observations (Cook's distance, residuals vs. leverage)

6. No multicollinearity (VIFs)

i Linear models in R

The following functions are useful when working with regression models in R:

- `lm` - fit the model
- `plot(lm)` - assess assumptions
- `fitted` - obtain est. values
- `resid` - residuals
- `predict` - forecast
- `model.matrix` - obtains design matrix (X)

Air Passengers activity

1. Load the `AirPassengers` data set and create a data frame that includes the count of air passengers, the year, the month, and a time index ($t = 1, 2, \dots, n$). Plot the count of air passengers by the time index. Then create a second data set, called `ap_sub`, that contains the air passenger measurements between 1949 and 1959.

! Important

Be sure to treat the month column as a **factor**!

2. Fit a regression model, called `fit`, that models the passenger count by the time index plus the month (be sure month is treated as a **factor**) for the `ap_sub` data frame. Print out a summary of the model, and use the summary to write out the estimated regression model.

3. Recreate the slope coefficient estimates by calculating the OLS estimator by hand using `model.matrix`. The following R functions may be useful:

- `solve` computes a matrix inverse
- `t` computes the transpose of a matrix
- `%*%` computes matrix multiplication

4. Interpret the slope coefficient associated with the time index.

5. Interpret the slope coefficient associated with the adjustment to the intercept for August.

6. Assess the linearity assumption for the regression model by running `plot(fit, which = 1)`

7. Assess the constant variance assumption for the regression model by using the previous plot and running `plot(fit, which = 3)`

8. Assess the normality assumption for the regression model by running `plot(fit, which = 2)`

9. Assess the independence assumption by thinking about the problem and running `acf(resid(fit))` and `pacf(resid(fit))`

10. One way to address the violations of the normality and constant variance assumptions is to log transform the response. To address the violation of the linearity assumption, it may be helpful to create a squared time index variable. Create new columns in the `ap_sub` data frame that log the number of passengers and computes the square of the time index. Then fit a new model, called `fit_log`, that models the logged passenger count by the time index, time index squared, and month.
11. Reassess the assumptions using the new model.

12. If the residuals are positively serially correlated, we will tend to underestimate the standard errors of the regression coefficients. What impact does this have when determining statistical significance of regression coefficients? It may be helpful to know that the t-test provided in the R output is calculated as

$$t = \frac{\hat{\beta}}{SE(\hat{\beta})}$$

13. With the exception of the independence assumption, hopefully you are convinced that the model created in 10 is a reasonable model for the `AirPassengers` data set. Next, calculate the fitted values by hand and compare them to the values from `fitted`.
14. Plot the observed and fitted time series on a single plot and comment on how well the model estimates the observed relationship.

15. Calculate the residuals by hand and compare them to the values obtained from `resid`.
16. Use the `fit_log` model to forecast the time series for the year of 1960 (which we had previously excluded) using the `predict` function. The `predict` function requires a `newdata` argument to obtain forecasts - this data frame must include all the predictors used in the model. Plot the observed, fitted, and forecasted series on a single plot and include a 95% prediction interval for the forecasted series. For more about predicting from `lm` models, run `?predict.lm`.