

Day 21 - ARIMA models

Introduction

We finally combine all all the stochastic error models we have considered this semester, culminating in the seasonal ARIMA model.

```
# packages
library(tidyverse)
library(lubridate)
library(forecast)
```

Review(ish)

Recall the random walk model, $x_t = x_{t-1} + w_t$, where w_t is a white noise series. Compute the *first-order differences* for the model: $\nabla x_t = x_t - x_{t-1}$.





Integrated processes

i Integrated model

A series $\{x_t\}$ is _____, $I(d)$, if the d^{th} difference of $\{x_t\}$ is white noise, denoted $\nabla^d x_t = w_t$. It can be shown that $\nabla^d \equiv (1 - B)^d$, where B is the backshift operator. Therefore, $\{x_t\}$ is $I(d)$ if

$$(1 - B)^d x_t = w_t$$

```
cbe <- read_delim("cbe.dat")
elec <- dplyr::select(cbe, elec)
elec_ts <- ts(elec$elec, start = 1958, freq = 12)
par(mfrow = c(2,2))
plot(elec_ts)
plot(diff(elec_ts))
plot(diff(log(elec_ts)))
plot(diff(diff(log(elec_ts))))
```



ARIMA models

i ARIMA models

A time series $\{x_t\}$ follows an $\text{ARIMA}(p, d, q)$ process if the d^{th} differences of the series are an $\text{ARMA}(p, q)$ process. That is

$$\theta_p(B)(1-B)^d x_t = \phi_q(B)w_t$$

You will not be expected to manipulate ARIMA models to determine stationarity or invertibility in this class, but it can be done.

```
(arima_fit <- arima(elec_ts, order = c(1, 1, 1)))
```

Call:


```
arima(x = elec_ts, order = c(1, 1, 1))
```

Coefficients:

	ar1	ma1
	-0.5577	0.4502
s.e.	0.1177	0.1194

sigma^2 estimated as 183803: log likelihood = -2954.51, aic = 5915.02

```
pacf(resid(arima_fit)) # not good
```



Seasonal ARIMA models

i Seasonal ARIMA models

Finally, we introduce the seasonal ARIMA model, in which we allow each of the autoregressive, integrated, and moving average components to depend on the value from the previous season. The seasonal ARIMA model, denoted $\text{ARIMA}(p, d, q)(P, D, Q)_s$, can be written in terms of the backshift operator:

$$\Theta_P(B^s)\theta_p(B)(1-B)^D(1-B)^d x_t = \Phi_Q(B^s)\phi_q(B)w_t$$

```
(arima_fit <- arima(elec_ts, order = c(1, 1, 1), seasonal = c(1, 1, 1)))
```

Call:


```
arima(x = elec_ts, order = c(1, 1, 1), seasonal = c(1, 1, 1))
```

Coefficients:

	ar1	ma1	sar1	sma1
	0.0908	-0.7409	-0.0224	-0.5318
s.e.	0.0856	0.0641	0.0813	0.0644

sigma^2 estimated as 22433: log likelihood = -2464.45, aic = 4938.89

```
pacf(resid(arima_fit)) # better, still not good
```



i Fitting ARIMA

The `auto.arima` function in the `forecast` package uses AIC to select the “best” seasonal ARIMA model. You may restrict to a subset of seasonal ARIMA models by setting `max.p`, `max.d`, `max.q`, `max.P`, `max.D`, `max.Q` equal to a non-negative integer. Additionally, the `auto.arima` function accepts covariates for regression via the `xreg` argument.

To forecast with an `auto.arima` fit, we use the `forecast` function. See `?forecast.Arima` for more details.

```
(auto_fit <- auto.arima(elec_ts))
```

Series: `elec_ts`

ARIMA(1,1,2)(0,1,1)[12]

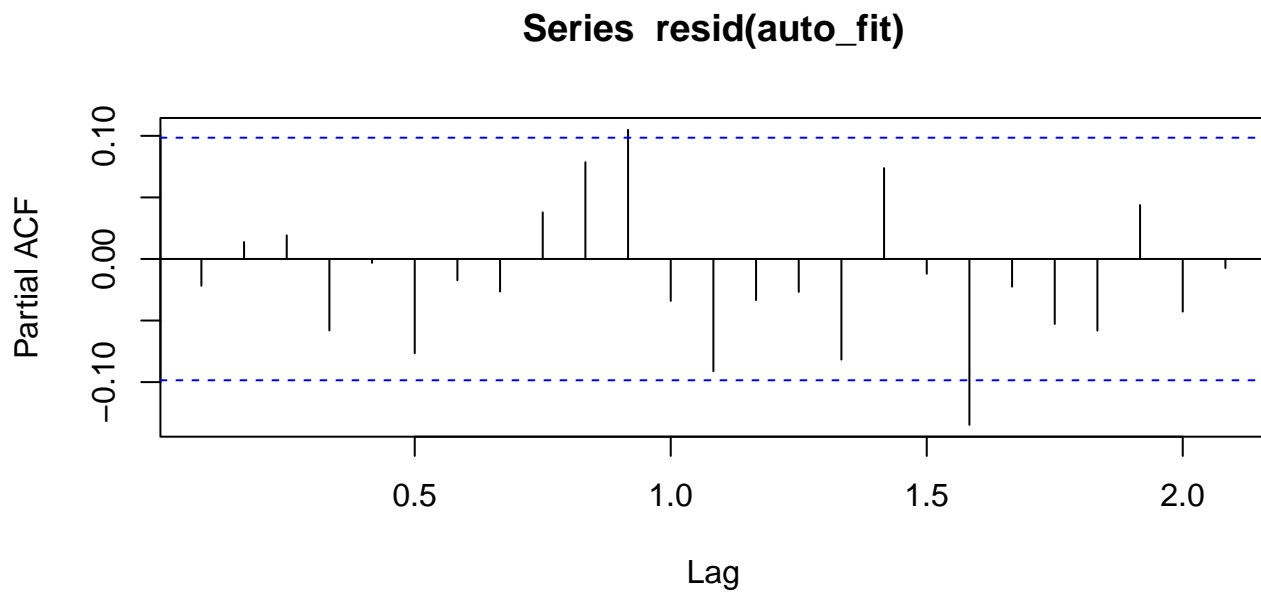
Coefficients:

	ar1	ma1	ma2	sma1
	0.8294	-1.4833	0.5125	-0.5368
s.e.	0.1063	0.1301	0.1096	0.0439

$\sigma^2 = 22390$: log likelihood = -2462.25

AIC=4934.51 AICc=4934.67 BIC=4954.25

```
pacf(resid(auto_fit))
```



```
# fitted is defined for auto.arima fits!
fitted(auto_fit) |> head()
```

	Jan	Feb	Mar	Apr	May	Jun
1958	1496.136	1462.640	1647.608	1594.754	1776.639	1823.657

```
# use the forecast function to forecast
auto_forecast <- forecast(auto_fit, h = 5*12, level = 95)
```

```
# putting it together
obs_df <- elec %>%
  mutate(time = c(time(elec_ts)), fitted = c(fitted(auto_fit))) %>%
  pivot_longer(-time)
```

```
forecast_df <- tibble(
  time = c(time(auto_forecast$mean)),
  value = c(auto_forecast$mean),
  lwr = c(auto_forecast$lower),
  upr = c(auto_forecast$upper)
) %>% mutate(name = "forecast")
```

```
ggplot() +
  geom_line(data = obs_df, aes(x = time, y = value, col = name)) +
  geom_line(data = forecast_df, aes(x = time, y = value, col = name)) +
  geom_ribbon(
    data = forecast_df, aes(x = time, ymin = lwr, ymax = upr),
    alpha = .30
  ) +
  theme_bw()
```