

## Homework 2

Be sure to submit **both** the .pdf and .qmd file to Canvas by Monday, September 23rd at 11:59 pm.

0. [1 pt] Did you work with anyone on this homework? If so, who?


1. [13 pt] For this homework, we continue practicing decomposition of time series, again recreating the output of `decompose`, but this time focusing on a multiplicative decomposition model.

- a) [1 pt] Load the `AirPassengers` data set.

```
data(AirPassengers)
ap <- AirPassengers
```

- b) [1 pt] Plot the raw `AirPassengers` time series and describe what you see in terms of trend and seasonal effect.

```
plot(ap)
```




The series has a strong positive trend and an equally strong seasonal component that increases in magnitude as the trend increases.

- c) [1 pt] Is an additive decomposition model appropriate for these data? Why or why not?

No. The magnitude of the seasonal effect increases as the trend increases, suggesting a multiplicative model would be more appropriate.

- d) [1 pt] Regardless of your answer to the previous question, fit a multiplicative decomposition model to the `AirPassengers` data set and plot the results.

```
ap_decompose <- decompose(ap, type = "multiplicative")  
plot(ap_decompose)
```



- e) [3 pt] Calculate the trend component for the multiplicative model by hand, and compare your estimated trend effects to those obtained by `decompose` to ensure they are correct.

```

ap_trend <- matrix(NA, nrow = length(ap) - 12, ncol = 1)
for(t in 7:(length(ap)-6)){
  ap_trend[t-6,] <- (
    .5*ap[t-6] + ap[t-5] + ap[t-4] + ap[t-3] +
    ap[t-2] + ap[t-1] + ap[t] + ap[t+1] +
    ap[t+2] + ap[t+3] + ap[t+4] + ap[t+5] + .5*ap[t+6]
  )*(1/12)
}

# append the NAs we skipped
ap_trend_full <- c(rep(NA, 6), ap_trend, rep(NA, 6))

# check if hand values are equal, up to numeric tolerance
all(abs(ap_trend_full - ap_decompose$trend) <= 1e-10, na.rm = T)

```

[1] TRUE

- f) [3 pt] Calculate the average seasonal effect associated with each month and center the resulting estimates. Compare these 12 values to those obtained from the `decompose` function to ensure you have calculated them correctly.

```

s_hat <- ap / ap_trend_full
s_bar_tmp <- colMeans(matrix(s_hat, ncol = 12, byrow = T), na.rm = T)
s_bar <- s_bar_tmp/mean(s_bar_tmp)

# check if hand values are equal, up to numeric tolerance
all(abs(s_bar - ap_decompose$figure) <= 1e-10)

```

[1] TRUE

- g) [3 pt] Calculate the residual error series from the previously created objects representing the trend and seasonal effects. Compare the by-hand calculation to the results from `decompose` to ensure that you have calculated the series correctly.

```

# calculate residual error series
res <- ap / (ap_trend_full * s_bar)


# check if hand values are equal, up to numeric tolerance
all(abs(res - ap_decompose$random) <= 1e-10, na.rm = T)

```

[1] TRUE

2. [5 pt] The moving average calculated by the `decompose` function is a type of **smoother** or **filter**. Smoothing/filtering algorithms generally seek to predict a response at time  $t$  based on observations from both before and after  $t$ , as we saw with the moving average in question 1. Another popular smoother is the **loess** smoother, which stands for locally estimated scatterplot smoothing, and can be performed in R with the `stl` function.
- a) [2 pt] The following code fits a loess smoother to the `AirPassengers` data and prints the trend estimate. How does this compare to the trend created by `decompose`? Do you notice any major differences?

```
loess_periodic <- stl(ap, s.window = "periodic")
plot(loess_periodic)
```



```
loess_periodic$time.series[,2]
```

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug
1949	127.1873	126.6495	126.1117	126.1989	126.2861	126.7330	127.1799	127.4162
1950	134.3390	135.1084	135.8777	135.9454	136.0131	137.2093	138.4055	141.3114
1951	159.7945	161.6914	163.5884	164.6017	165.6150	167.2285	168.8420	171.2070
1952	185.7223	187.8163	189.9103	191.1312	192.3520	194.0658	195.7796	198.7814
1953	218.0156	220.1263	222.2371	222.9376	223.6380	223.7928	223.9476	223.9443
1954	228.7944	230.5535	232.3126	233.8328	235.3531	237.7041	240.0551	243.2465
1955	262.3979	266.7545	271.1111	274.8384	278.5657	282.0862	285.6067	289.3515
1956	309.7570	314.0295	318.3020	321.7071	325.1122	327.6514	330.1906	332.4831
1957	347.2005	351.9113	356.6220	360.9900	365.3580	368.2253	371.0926	372.2017


1958	373.5057	375.3037	377.1017	379.1766	381.2514	382.8734	384.4954	386.4129
1959	399.6645	404.9035	410.1425	416.3450	422.5474	427.7897	433.0319	436.4362
1960	452.3700	457.8121	463.2543	467.6104	471.9665	475.6083	479.2501	483.1894
	Sep	Oct	Nov	Dec				
1949	127.6525	129.0186	130.3846	132.3618				
1950	144.2173	148.3371	152.4568	156.1257				
1951	173.5721	176.5526	179.5332	182.6278				
1952	201.7832	206.0687	210.3542	214.1849				
1953	223.9409	224.5947	225.2486	227.0215				
1954	246.4380	250.1037	253.7693	258.0836				
1955	293.0963	297.0760	301.0557	305.4064				
1956	334.7756	337.3493	339.9230	343.5618				
1957	373.3108	373.0689	372.8270	373.1663				
1958	388.3304	390.6746	393.0187	396.3416				
1959	439.8406	442.3755	444.9104	448.6402				
1960	487.1288	490.7530	494.3773	497.4299				

The values differ by a little, and you are able to obtain trend estimates for the first and last 6 months.


- b) [2 pt] The code above assumes the window for calculating the loess smoother is based on the period of the time series. Play around with changing the `s.window` argument (which can be an integer) in the `stl` function. What do you notice about the resulting decomposition?

Changing the values can have a fairly significant impact on the magnitude of the estimate of the seasonal component.


```
plot(stl(ap, s.window = 3))
```



```
plot(stl(ap, s.window = 10))
```



```
plot(stl(ap, s.window = 30))
```



- c) [1 pt] The smoothing and filtering algorithms provide a nice way to summarize a time series in retrospect. Can you think of a shortcoming for using these tools to estimate time series and create forecasts?

You cannot create forecasts with them! They rely on data from time points both before and after the point of prediction, meaning you cannot forecast into the future with them. Generally, they do not provide a model that you can write down to perform forecasts.

3. [6 pt] Find an interesting **discrete** time series data set and create a decomposition of it below. Note that [Kaggle](#) and [Data.world](#) often provide some nicely manicured, free data.